

# **Harnessing The Power of Crossplane and Dapr**

Proof of Concept:

A Kyverno Policy Validator for Crossplane  
Compositions using Crossplane Functions and Dapr

[Hugo Smitter](#)

Platform Architect

[hugosmitter@fico.com](mailto:hugosmitter@fico.com)

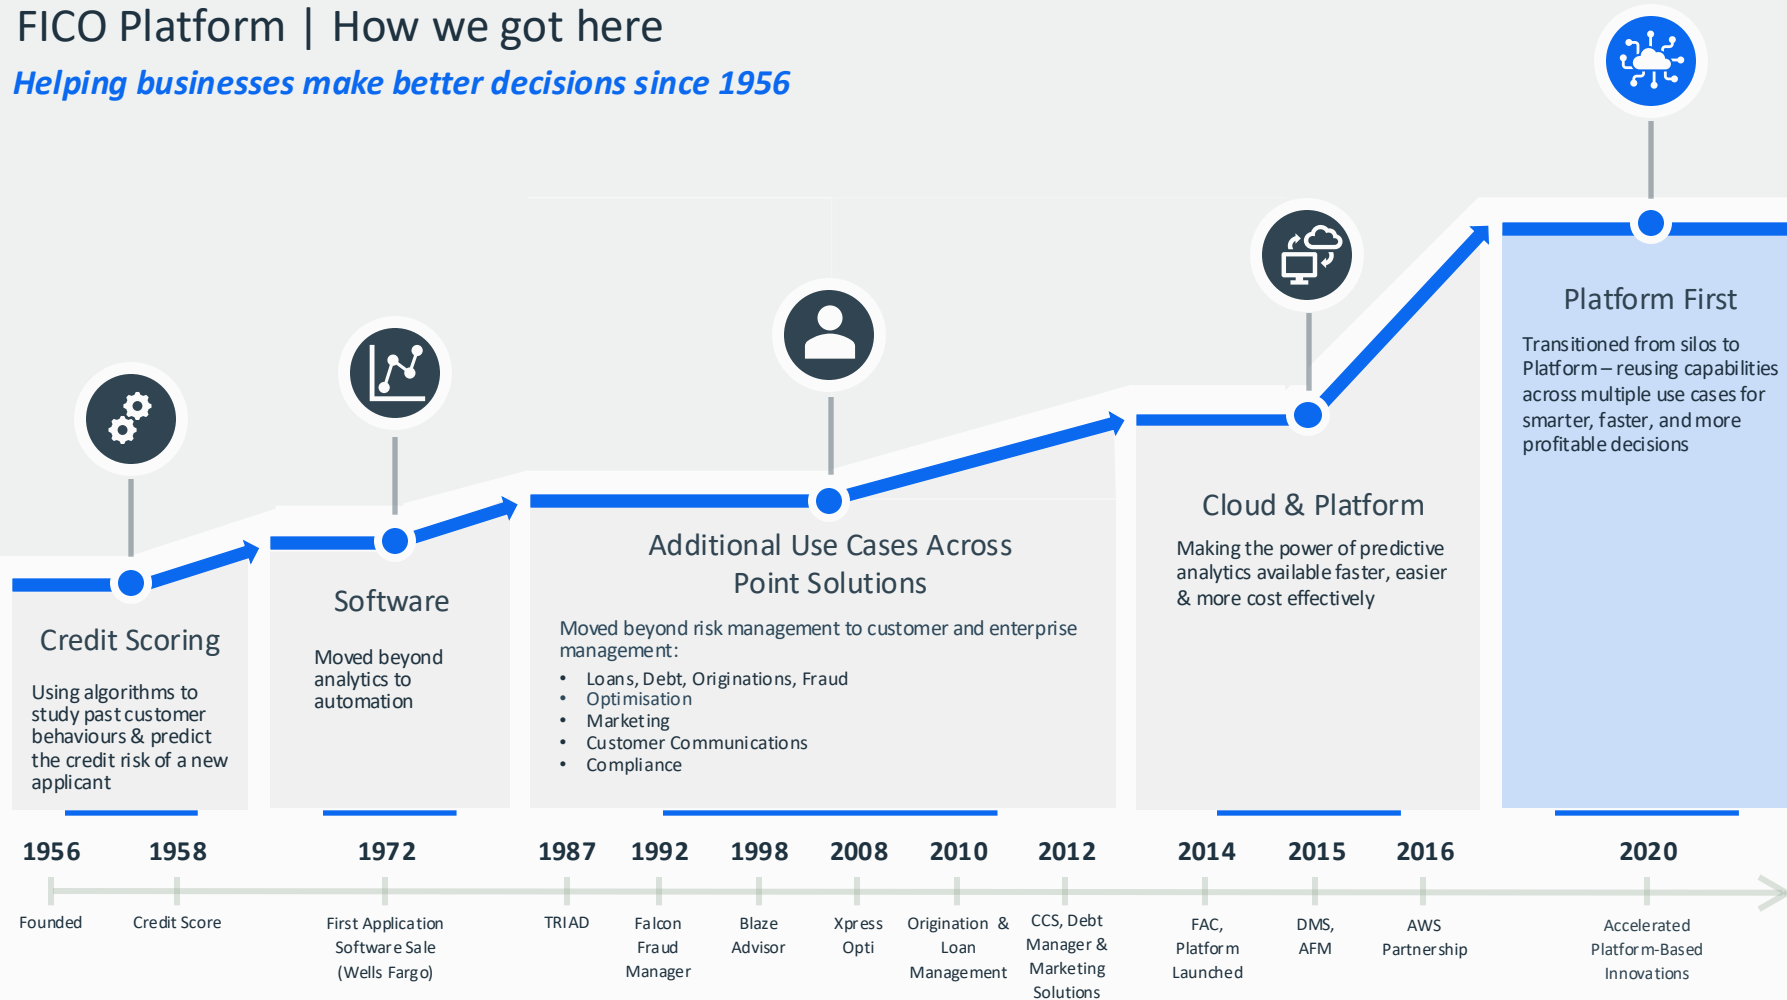
(+1) 214 534 5373

# Who is FICO and what do we do?



# FICO Platform | How we got here

*Helping businesses make better decisions since 1956*



# Agenda

- Crossplane overview
- Opportunities and challenges when using Crossplane (Composition) Functions
- Dapr overview: Build more complex but maintainable Crossplane Functions
- How to integrate Dapr and Crossplane
- Proof-of-Concept - A Kyverno Policy Validator for Crossplane Compositions:
  - Problem we're trying to solve
  - Challenges and solutions
  - Architecture diagrams
- Demo
- Next steps
- Q&A
- Appendix



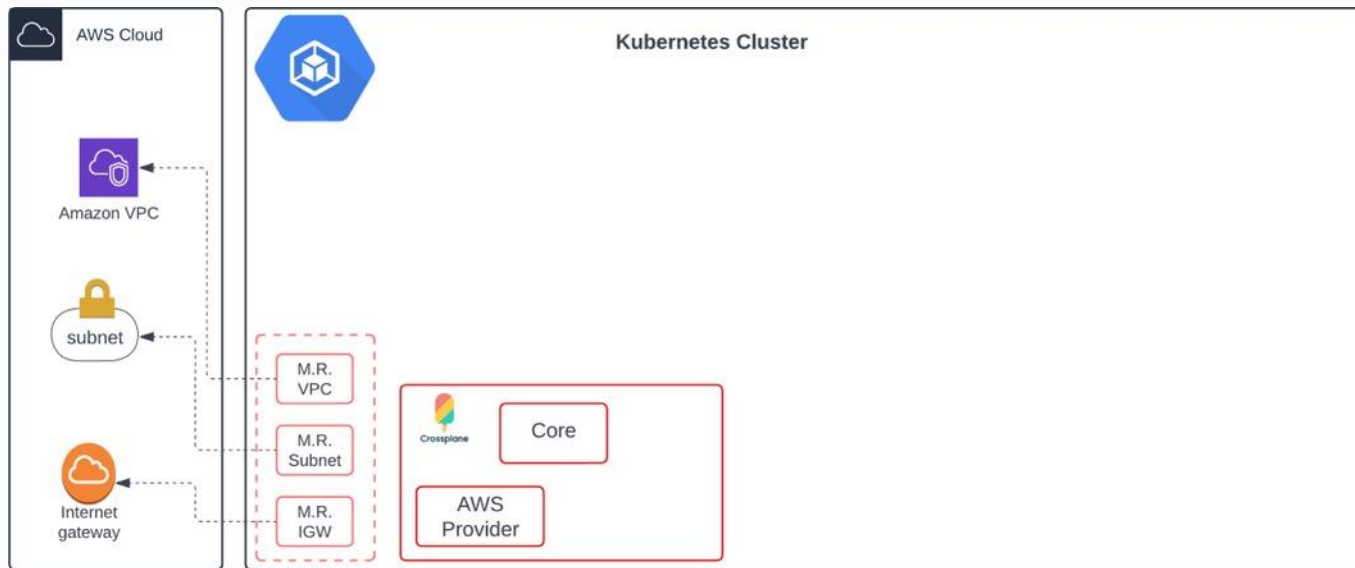
# Quick Crossplane Overview




# Crossplane Overview

Managed Resources (M.R.) are representations of external resources in Crossplane

Providers enable Crossplane to provision things (e.g.: infrastructure on an external service). They create new K8s APIs and map them to external APIs (\*).



## Kubernetes Control Loop



(\*) <https://marketplace.upbound.io/providers>

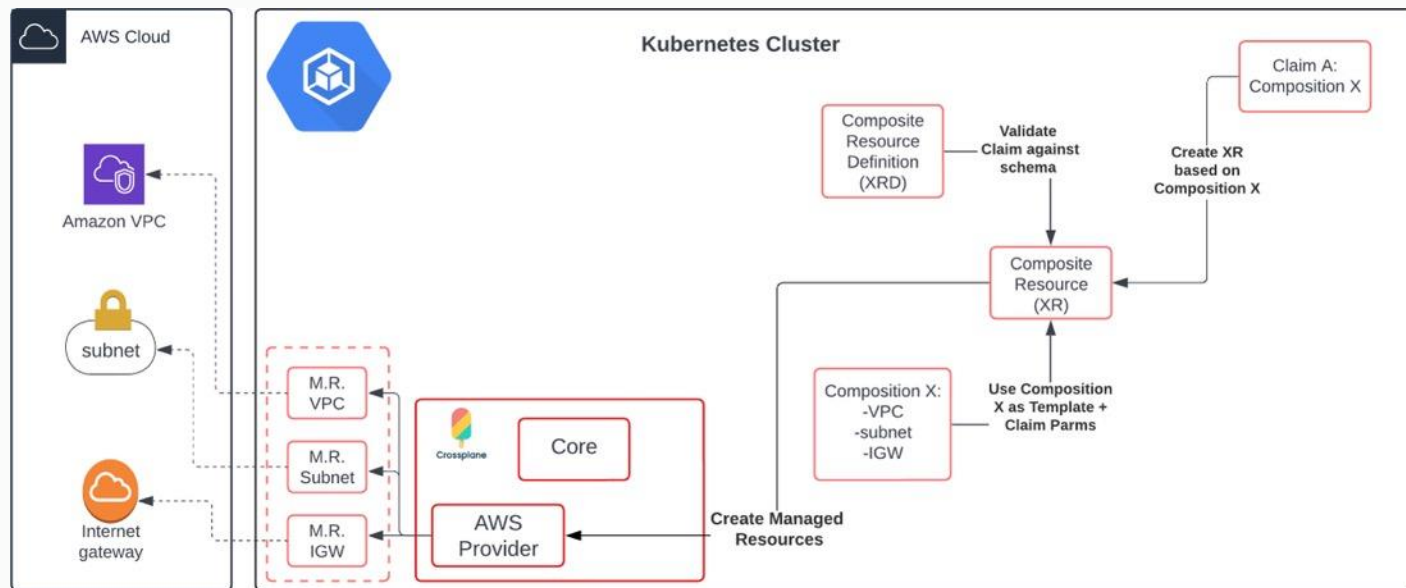
## Crossplane Overview (cont.)

Claims represent a set of managed resources as a single K8s object *inside a namespace*

Composite Resource Definitions (XRDs) define the schema for a custom API (claim)

Compositions are templates for creating multiple Managed Resources (MR)

Composite Resources (XRs) represent a set of MRs as a single K8s object

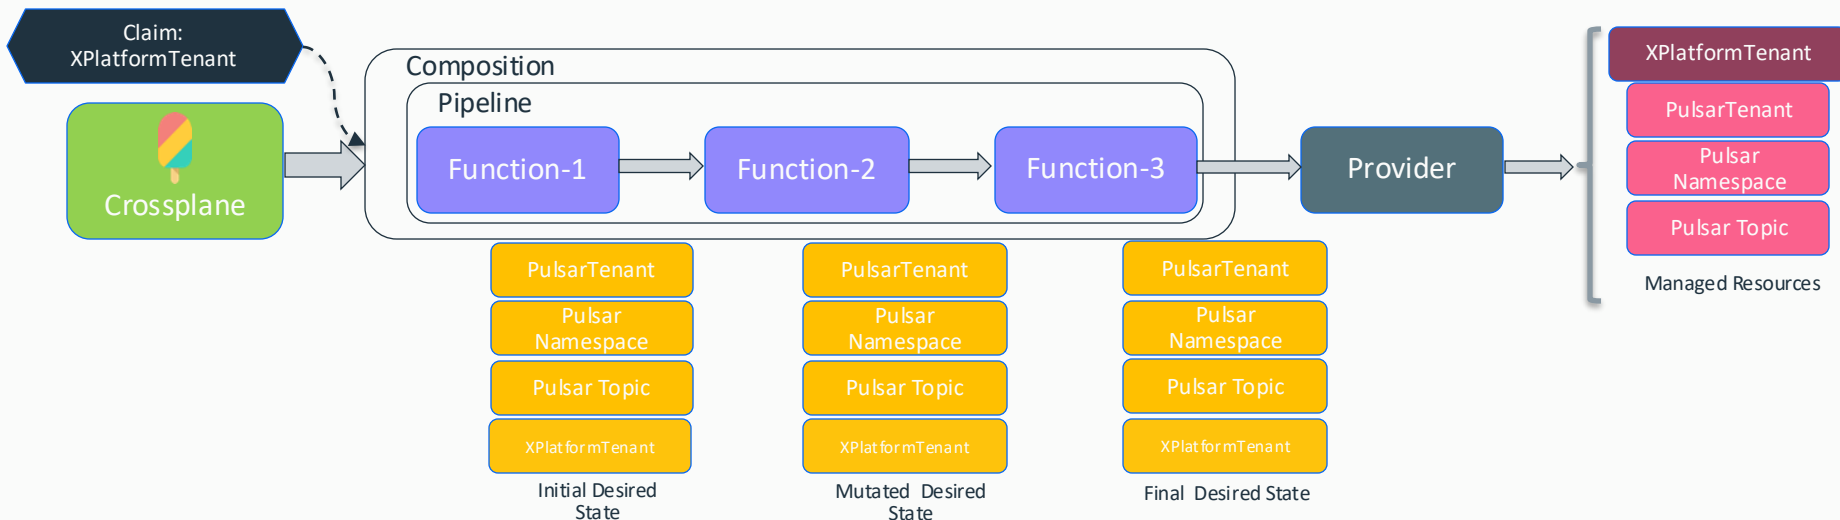


## Crossplane Overview (cont.)

<https://marketplace.upbound.io/functions?query=functions>

Composition Functions are custom programs that template Crossplane resources.

- Functions dynamically determine what resources should be created when you create composite resource (XR).
- You can write Functions in Go or Python (more languages to come).
- You can write advanced logic to template resources, like loops and conditionals. You can also access data or call other programs.
- Compositions can have pipelines of functions assembled in steps, where the output of a function is the input to the next function.



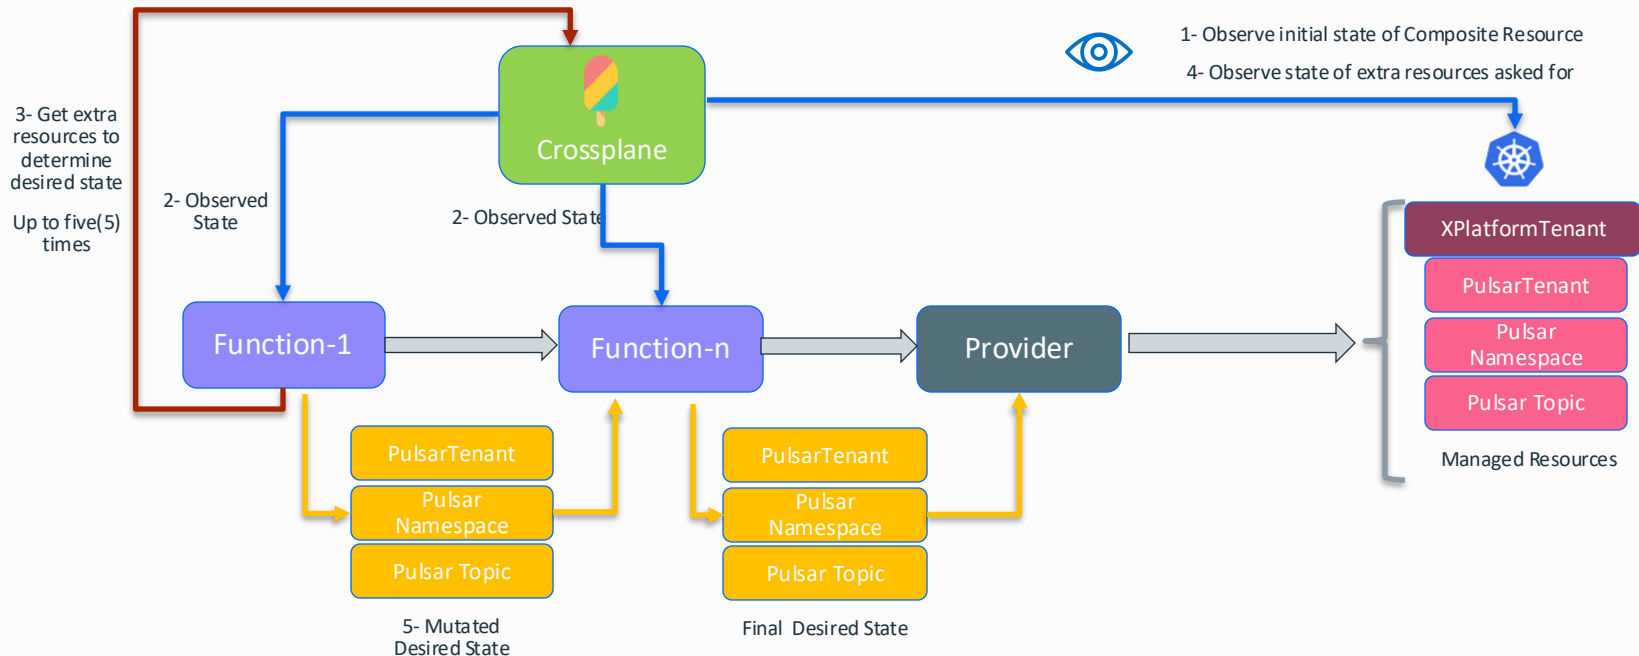
<https://www.youtube.com/watch?v=n0j5WiavXEI> = Steven Borrelli – Upbound Inc.



# Crossplane Overview (cont.)

## Empirical observations on how compositions and functions work:

- They work in loops! Their behavior is more complex than a sequence of steps in a pipeline (\*)
  - Composed/Managed Resource **provisioning loop**: Five (5) iterations until stability is reached or error
  - **Reconciliation loop** once resources are created (Default: every minute)



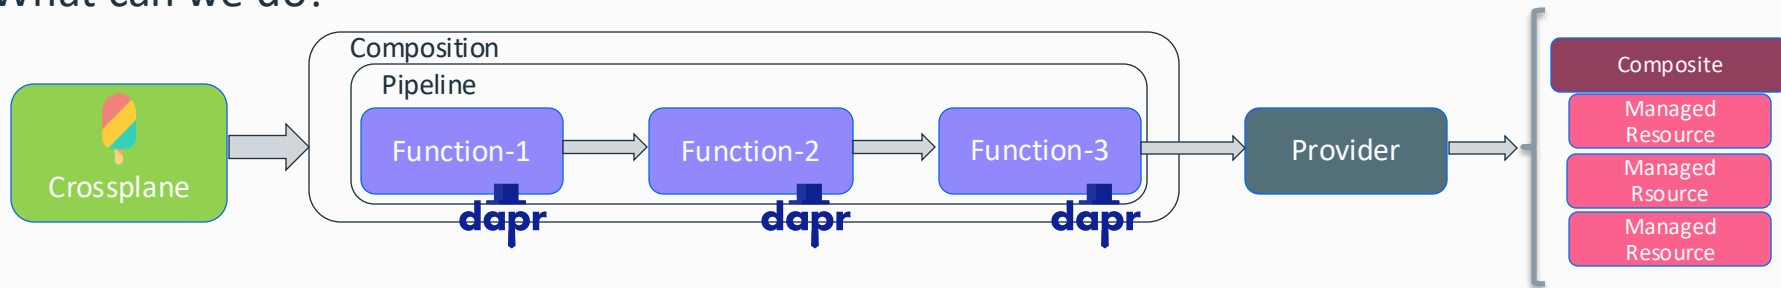
(\*) <https://docs.crossplane.io/latest/concepts/compositions/#how-composition-functions-work>

# Some Challenges Provisioning Anything-As-Code Even With Functions




- Coding more advanced logic in Functions with good separation of concerns
- Dealing with those provisioning and reconciliation loops!
- Using state stores, secrets vaults, pub/sub middleware, workflows, encryption and other complex technologies without further increasing composition/pipeline/function complexity
- Writing to common APIs for external services running in different technology stacks, cloud providers and/or using multiple programming languages
- Doing the above while developing and testing locally, followed by a seamless deployment in Kubernetes

## What can we do?



Quick  
**dapr**  
Overview



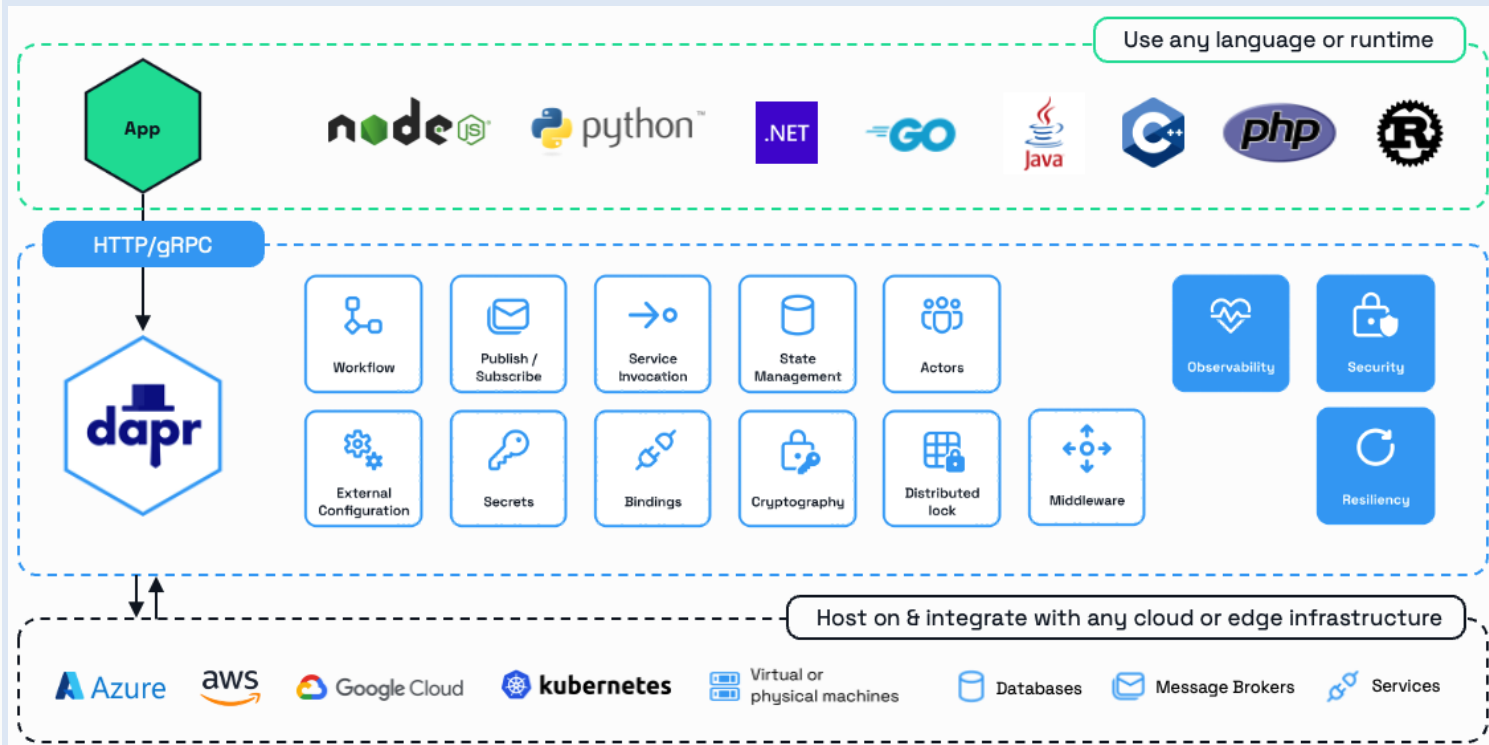


Distributed Application Runtime

Portable, event-driven, runtime for building distributed applications across cloud and edge

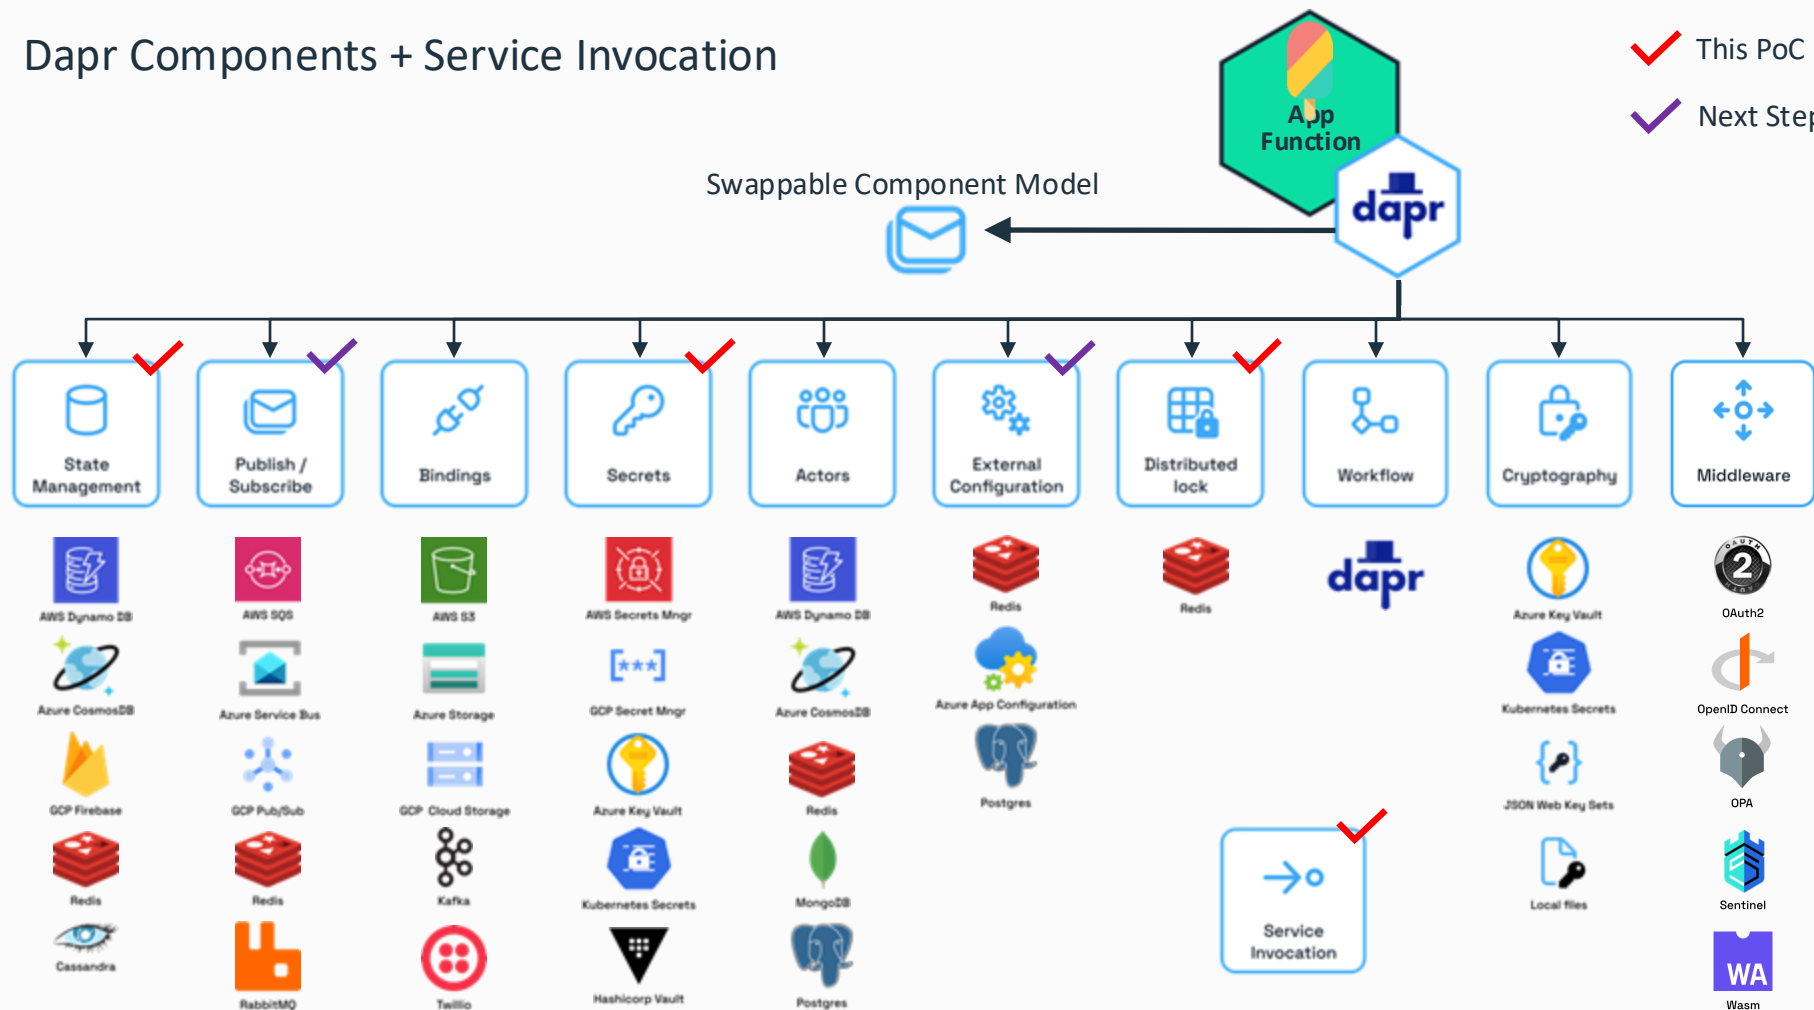
<https://dapr.io>

<https://docs.dapr.io/contributing/presentations/>

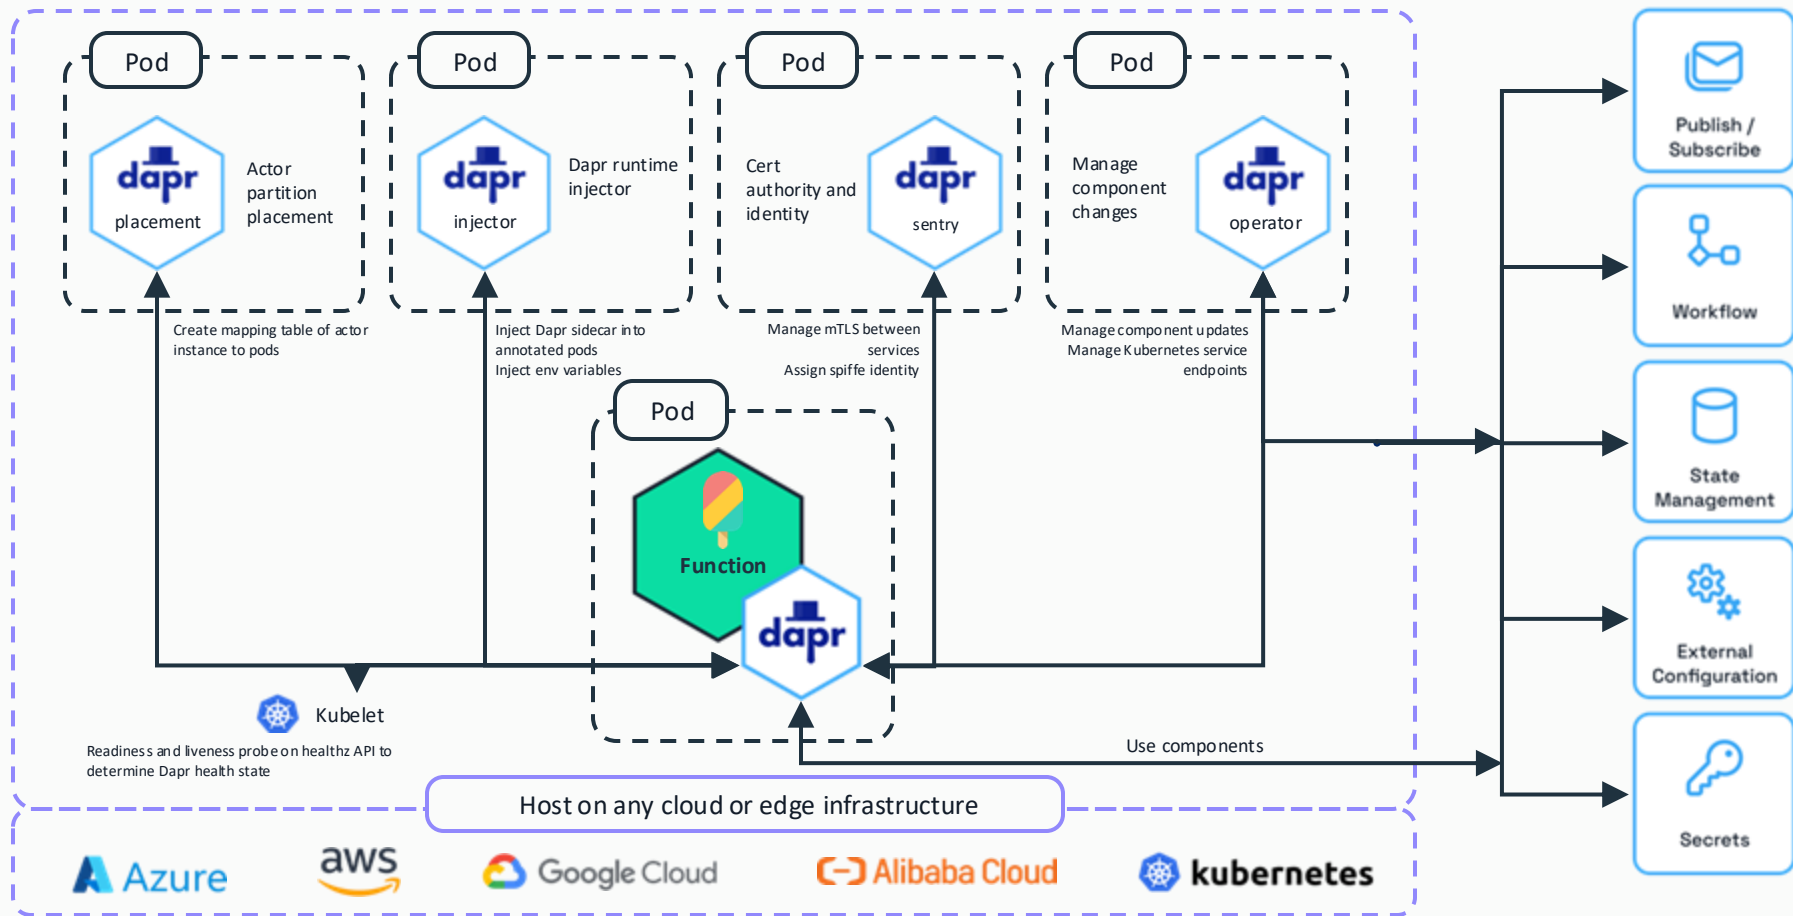


# Dapr Components + Service Invocation

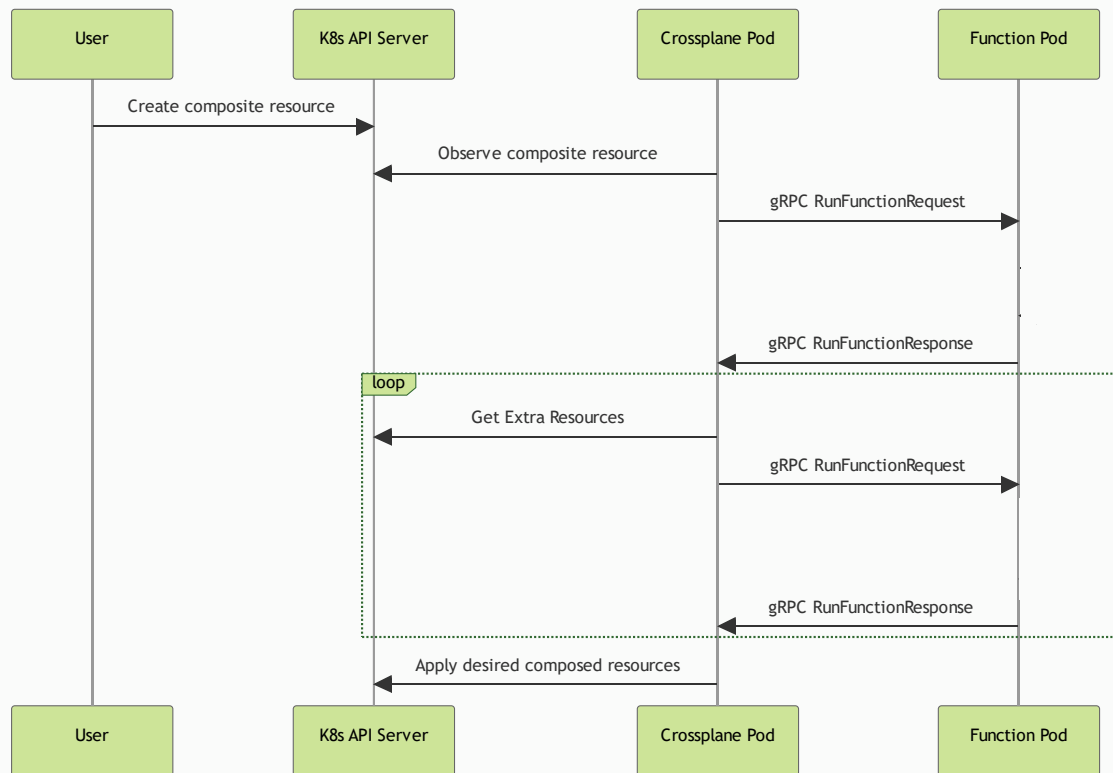
✓ This PoC  
✓ Next Steps



# Dapr on Kubernetes



# How Composition Functions work and how do the work with Dapr



<https://docs.crossplane.io/latest/concepts/compositions/#how-composition-functions-work>

## Prerequisite steps for using Dapr

- Install the `dapr cli` in your laptop and run a `dapr init` command
- Install Dapr in your Kubernetes cluster
  - `dapr init -k` or
  - Run the Dapr installation Helm chart for greater control over configuration:  
<https://docs.dapr.io/operations/hosting/kubernetes/kubernetes-deploy/>



# Minimum modifications to a Composition Function template project in Python

The screenshot shows the Crossplane documentation website. The browser address bar displays `docs.crossplane.io/latest/guides/write-a-composition-function-in-python/`. The page header includes the Crossplane logo, navigation links for 'Why Control Planes?', 'Documentation', 'Community', and 'Blog', and a version selector set to 'v1.17 Latest'. A left sidebar contains a search bar and a list of navigation items: Overview, Getting Started, Install, Upgrade and Uninstall, Concepts, Guides (selected), Disaster Recovery with Crossplane, Metrics, Function Patch and Transform, Write a Composition Function in Go, Write a Composition Function in Python (highlighted), Import Existing Resources, Vault as an External Secret Store, Vault Credential Injection, Multi-Tenant Crossplane, Configuring Crossplane with Argo CD, Self-Signed CA Certs, Troubleshoot Crossplane, CLI Reference, API Reference, Learn More, Contributing Guide, and Crossplane Roadmap. The main content area is titled 'Write a Composition Function in Python'. It contains an introductory paragraph about composition functions, a code sample for an `XBuckets` resource, and a list of steps to follow. An 'Important' callout box is also present.

docs.crossplane.io/latest/guides/write-a-composition-function-in-python/

Crossplane

Why Control Planes? Documentation Community Blog

Search

Overview

Getting Started

Install, Upgrade and Uninstall

Concepts

Guides

Disaster Recovery with Crossplane

Metrics

Function Patch and Transform

Write a Composition Function in Go

Write a Composition Function in Python

Import Existing Resources

Vault as an External Secret Store

Vault Credential Injection

Multi-Tenant Crossplane

Configuring Crossplane with Argo CD

Self-Signed CA Certs

Troubleshoot Crossplane

CLI Reference

API Reference

Learn More

Contributing Guide

Crossplane Roadmap

## Write a Composition Function in Python

Composition functions (or just functions, for short) are custom programs that template Crossplane resources. Crossplane calls composition functions to determine what resources it should create when you create a composite resource (XR). Read the [concepts](#) page to learn more about composition functions.

You can write a function to template resources using a general purpose programming language. Using a general purpose programming language allows a function to use advanced logic to template resources, like loops and conditionals. This guide explains how to write a composition function in [Python](#).

**Important**

It helps to be familiar with [how composition functions work](#) before following this guide.

### Understand the steps

This guide covers writing a composition function for an `XBuckets` composite resource (XR).

```
1  apiVersion: example.crossplane.io/v1
2  kind: XBuckets
3  metadata:
4    name: example-buckets
5  spec:
6    region: us-east-2
7    names:
8      - crossplane-functions-example-a
9      - crossplane-functions-example-b
10     - crossplane-functions-example-c
```

An `XBuckets` XR has a region and an array of bucket names. The function will create an Amazon Web Services (AWS) S3 bucket for each entry in the names array.

To write a function in Python:

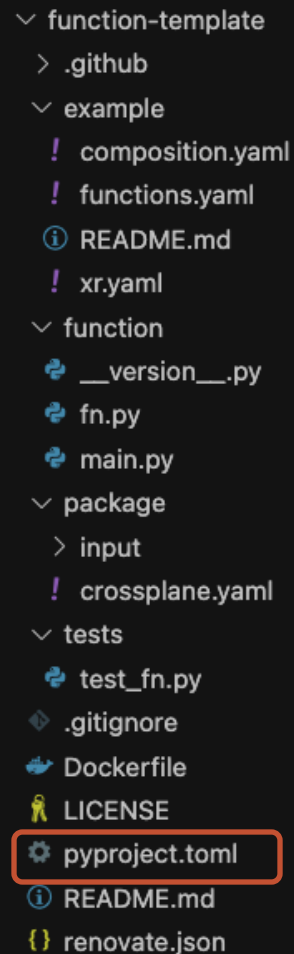
1. Install the tools you need to write the function

## Minimum modifications to a Composition Function template project in Python

- Add Python `requests` package to `pyproject.toml`
- That's all you need to call Dapr via HTTP from a Crossplane Function

Add a current  
version of the  
Python  
'requests'  
package

```
dependencies = [  
    → "requests==2.32.3",  
    "crossplane-function-sdk-python==0.5.0",  
    "click==8.1.7",  
    "grpcio==1.67.0",  
]
```



## Minimum modifications to a Composition Function template (any language)

- Add a DeploymentRuntimeConfig manifest to your project

Add name  
to function  
Config Ref  
parameter

```
# DeploymentRuntimeConfig to inject Dapr to a function
# Change dapr.io/app-id annotation to match your function name
apiVersion: pkg.crossplane.io/v1beta1
kind: DeploymentRuntimeConfig
metadata:
  name: function-policy-validation-deployment-config
spec:
  deploymentTemplate:
    metadata:
      labels:
    spec:
      selector: {}
      template:
        metadata:
          annotations:
            dapr.io/enabled: "true"
            dapr.io/scheduler-host-address: ""
            dapr.io/app-id: "<your-function-name>"
```

Inject the  
Dapr sidecar  
when  
deploying  
the function

```
function-template
├── .github
├── example
│   └── ! composition.yaml
│       └── ! functiondeploymentconfig.yaml
│           └── ! functions.yaml
│               └── README.md
│                   └── ! xr.yaml
├── function
│   ├── __version__.py
│   ├── fn.py
│   └── main.py
├── package
│   └── input
│       └── ! crossplane.yaml
├── tests
│   └── test_fn.py
├── .gitignore
├── Dockerfile
├── LICENSE
├── pyproject.toml
├── README.md
└── renovate.json
```

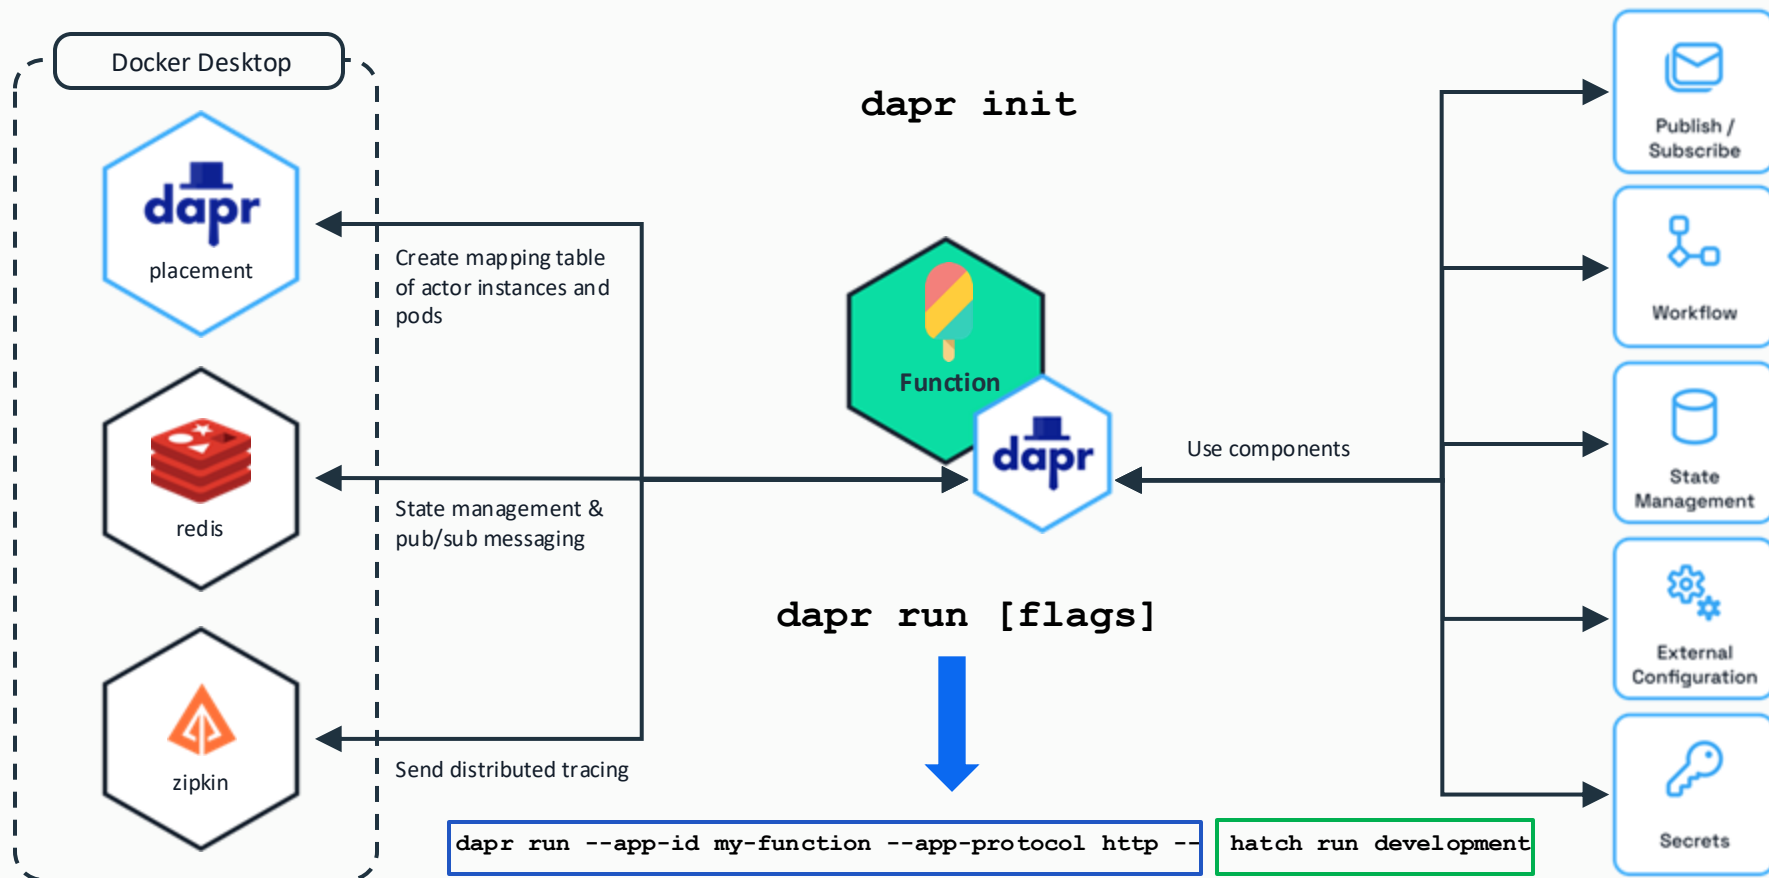
## Minimum modifications to a Composition Function template (any language)

- Add reference to your DeploymentRuntimeConfig manifest for your functions.yaml manifest

```
apiVersion: pkg.crossplane.io/v1
kind: Function
metadata:
  name: function-policy-validation
  annotations:
    # This tells crossplane beta render to connect to the function locally.
    render.crossplane.io/runtime: Development
spec:
  # This is ignored when using the Development runtime.
  package: <function docker image name>:<tag>
  runtimeConfigRef:
    { name: function-policy-validation-deployment-config
```

```
function-template
├── .github
├── example
│   ├── ! composition.yaml
│   ├── ! functiondeploymentconfig.yaml
│   └── ! functions.yaml
├── ! README.md
├── ! xr.yaml
├── function
│   ├── __version__.py
│   ├── fn.py
│   └── main.py
├── package
│   └── input
│       └── ! crossplane.yaml
├── tests
│   └── test_fn.py
├── .gitignore
├── Dockerfile
├── LICENSE
├── pyproject.toml
├── ! README.md
└── {} renovate.json
```

# Local development with the Dapr CLI



Putting it all together



+





# PoC: Kyverno Policy Validator for Crossplane



Kyverno





# Problem we're trying to solve in this PoC



## Current Status:

1. The Kyverno project provides a comprehensive set of tools to manage the complete Policy-as-Code (PaC) for Kubernetes and other cloud native environments.
2. Kyverno CLI can be used to apply and test policies off-cluster e.g. in IaC and CI/CD pipelines.
3. Policy enforcement happens at the “last mile” in Kubernetes admission controllers.
4. We want to make broad use of Kyverno policies for “Everything-as-Code”.
5. We don't want to mutate in the admission controllers and clash with Crossplane.

## Gap:

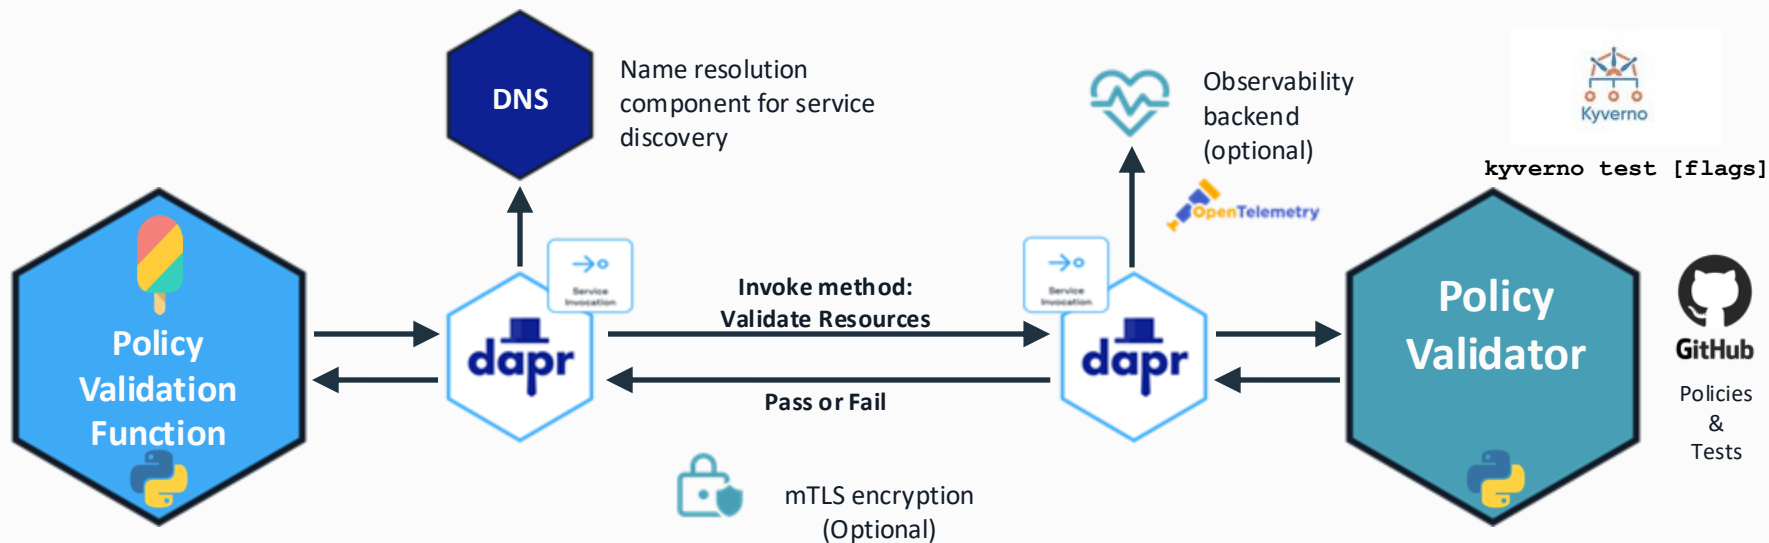
- With Crossplane Functions, resources could be generated **dynamically**:
  - Bullet #2 Happens too early. Not all resources to be generated may exist.
  - Bullet #3 Happens too late to validate Compositions prior to deployment

## Proposed Solution:

- A Policy Validation Function to be called towards the end of a Composition. It will wrap a **kyverno test** command against the top Composite resource generated in-flight and the applicable policies and tests for the top Composite resource stored in a GitHub repo.



# Policy Validation Function and PolicyValidator Microservice



POST

<http://localhost:3500/v1.0/invoke/policyvalidator/method/validate>

POST

<http://localhost:5100/pass>

POST

<http://localhost:5100/fail>

POST

<http://localhost:5100/locked>

# Some of PoC challenges and solutions

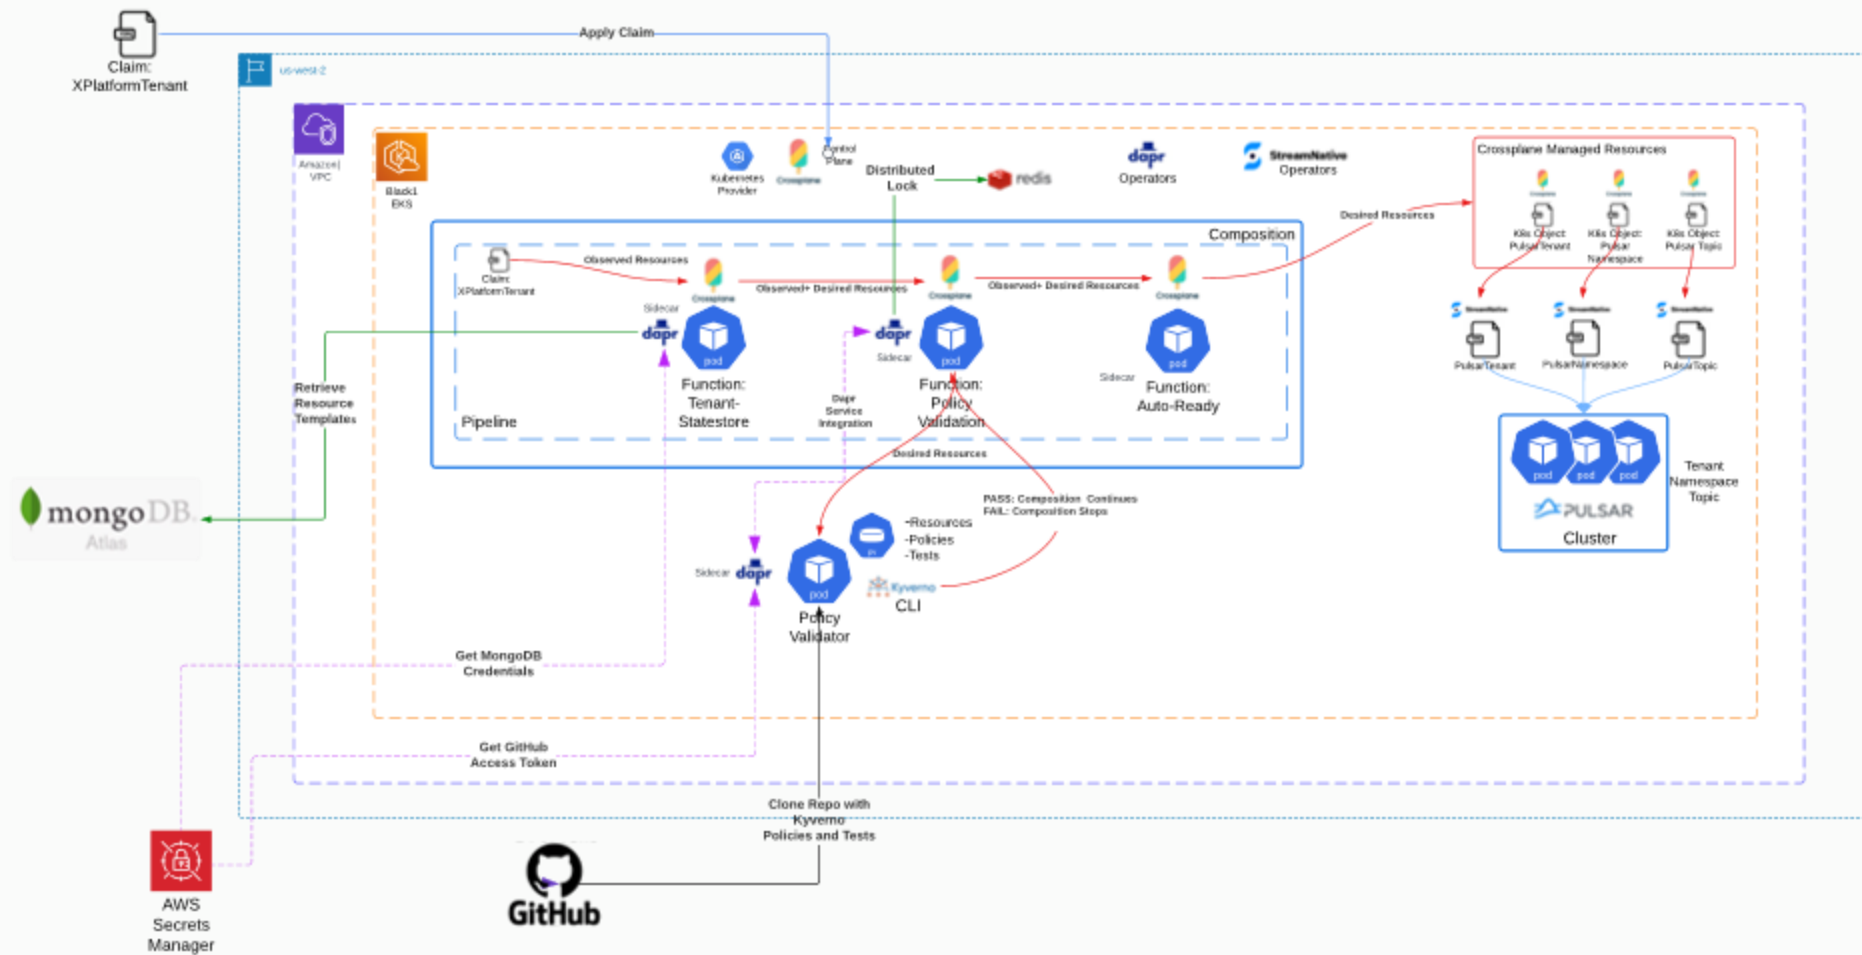


## Challenges: Complexity writing a monolithic function! The mission:

- a) Gather all the desired state resources and render in YAML for the kyverno CLI
- b) Securely `git clone` the repo with the applicable policies and tests
- c) Store the data in a file system in the right format and structure
- d) Call the `kyverno test` command, format the output, act based on `pass/fail` and clean up file system
- e) Handle the Crossplane the creation loop mechanics and don't break the Composition

## Solutions:

- A function to call a PolicyValidator microservice using Dapr Service Invocation (HTTP). Offload the heavy lifting
- Handle transformations in Python: Crossplane proto → JSON strings → Python dict → YAML (Booleans are tricky)
- Dapr Secrets Management to retrieve access token to policies/tests GitHub repo
- Dapr Distributed Lock to lock PolicyValidator (PV) microservice for X minutes after first call until it can return a `pass`, `fail` or `locked` back to the calling function. Three outcomes:
  1. PV returns `pass`: Calling function returns desired state resources unaltered from the prior function.
  2. PV returns `fail`: Calling function fails the composition returning an abnormal termination log entry.
  3. PV returns `locked`: Similar to a `pass`, but with a warning log entry to look for `pass` or `fail` log entry
- PolicyValidator file system and repo naming convention: `kind-apiVersion`. Concurrency out of scope for PoC.



## Next Steps



- Create a function to update a state store and send an event to an event service using the Transactional Outbox pattern (\*). Make reporting Kyverno cli test results a bit friendlier
- Test the Dapr Configuration building block to facilitate configuring function parameters
- Upload project: functions, scripts, configuration file, etc. to a GitHub repository
- Submit PR to the Dapr community and Diagrid to add a “create” method to the Secrets Management building block to create new secrets.
  - This is not typical application developer function, but a necessary one for platform engineering
- Socialize Crossplane/Dapr integration with Upbound and the Crossplane community

(\*) <https://docs.dapr.io/developing-applications/building-blocks/state-management/howto-outbox/>

# Thank You!

Hugo Smitter

Platform Architect

[hugosmitter@fico.com](mailto:hugosmitter@fico.com)

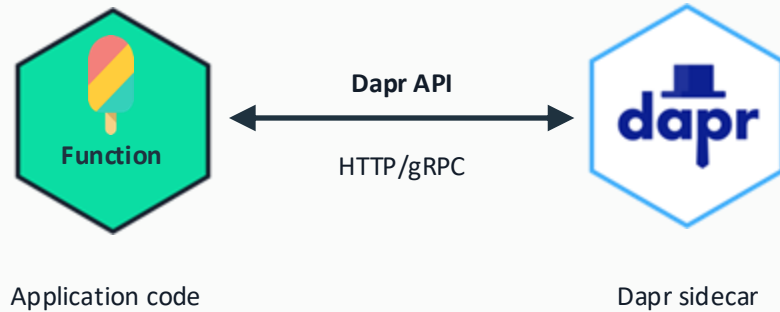
(+1) 214 534 5373



<http://github.com/smitterh/Crossplane-Dapr-Integration>

# Appendix

## Sidecar pattern and the Dapr API



**POST** <http://localhost:3500/v1.0/invoke/cart/method/order>


**GET** <http://localhost:3500/v1.0/state/inventory/item50>

**POST** <http://localhost:3500/v1.0/publish/mybroker/order-messages>

**GET** <http://localhost:3500/v1.0/secrets/vault/dbaccess>

**POST** <http://localhost:3500/v1.0-beta1/workflows/dapr/businessprocess/start>

# State Management (Key/Value)




key	field	value
template-pulsar-tenant	data	"{apiVersion: ...}"
template-pulsar-namespace	data	"{apiVersion: ...}"
template-pulsar-topic	data	"{apiVersion: ...}"

GET

<http://localhost:3500/v1.0/state/my-statestore/template-pulsar-tenant>



## State Management (Key/Value) (cont.)





key	field	value
template-pulsar-tenant	data	"{apiVersion: ...}"
template-pulsar-namespace	data	"{apiVersion: ...}"
template-pulsar-topic	data	"{apiVersion: ...}"

**POST**

<http://localhost:3500/v1.0/state/my-statestore>

```
{
  "key": "template-pulsar-tenant",
  "value": "{apiVersion: x, kind: PulsarTenant, metadata: ...}"
}
```

## Secrets Management (use case 1)



**GET** <http://localhost:3500/v1.0/secrets/myvault/mysecret>

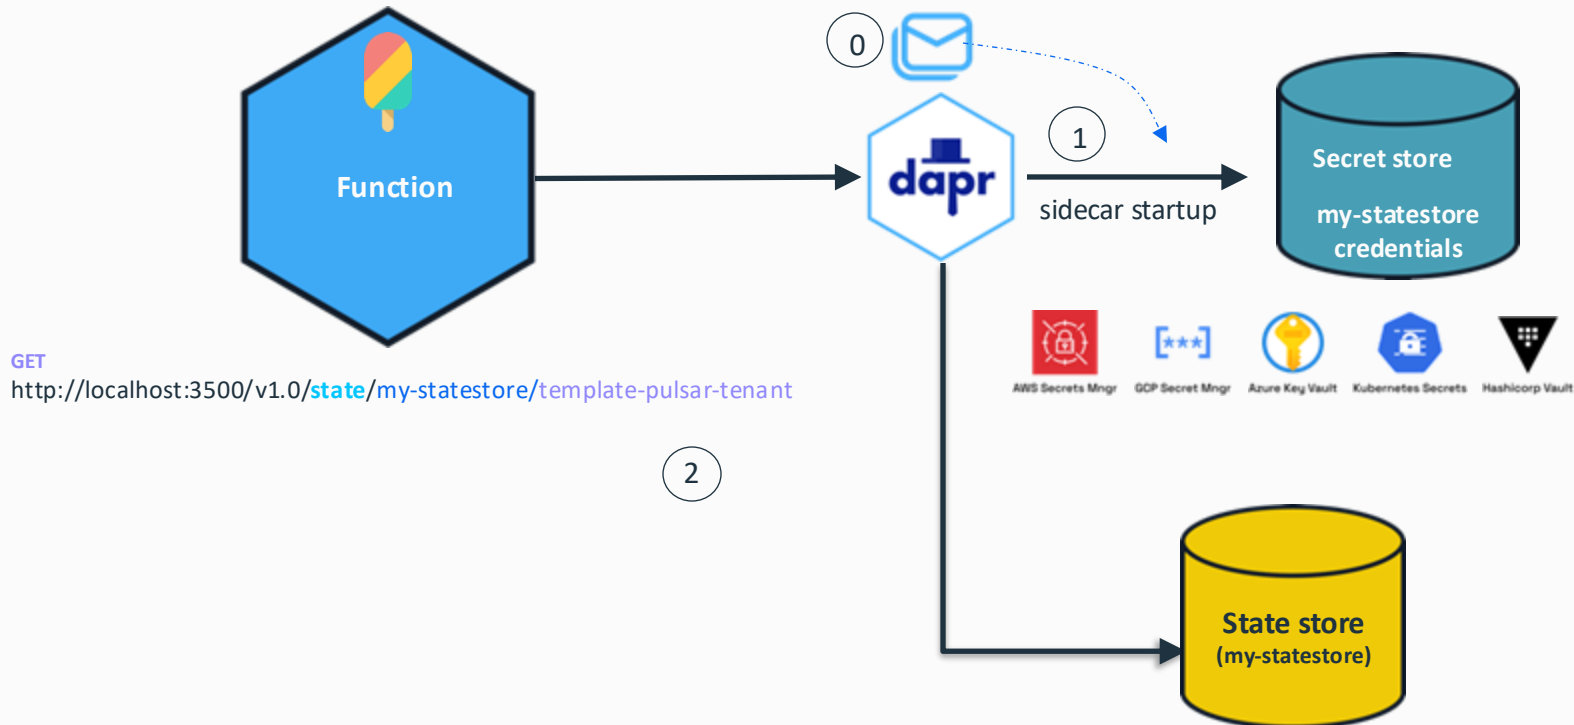
RESPONSE

```
{
  "mysecret": "secretvalue"
}
```

## Secrets Management (use case 2)



Configure State Store Component Manifest



# Distributed Lock

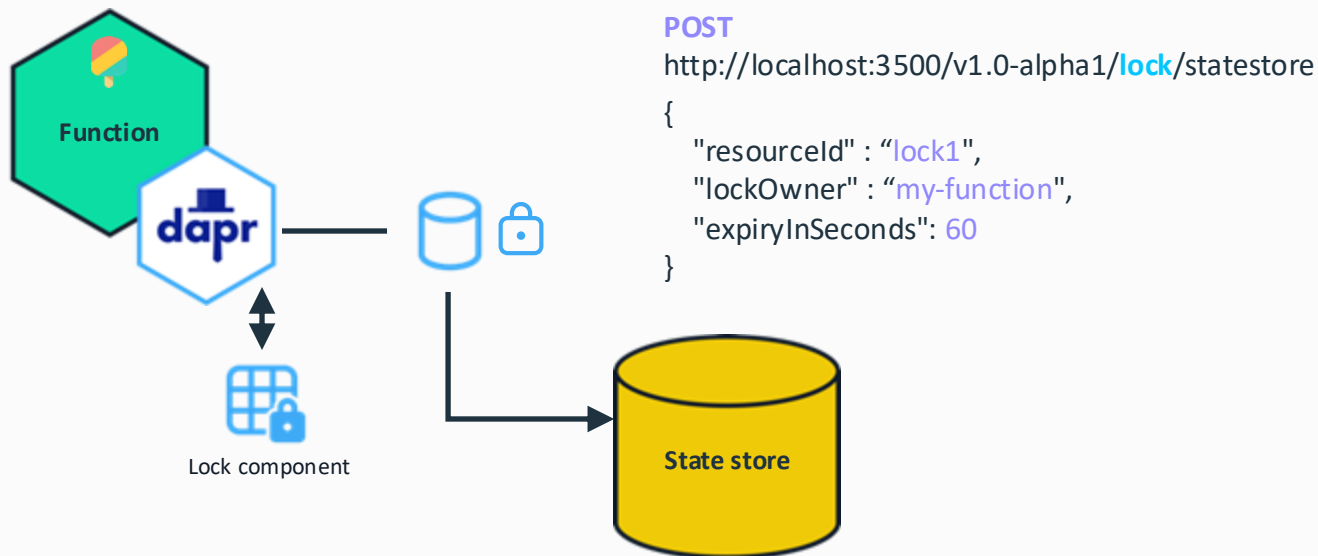


ALPHA

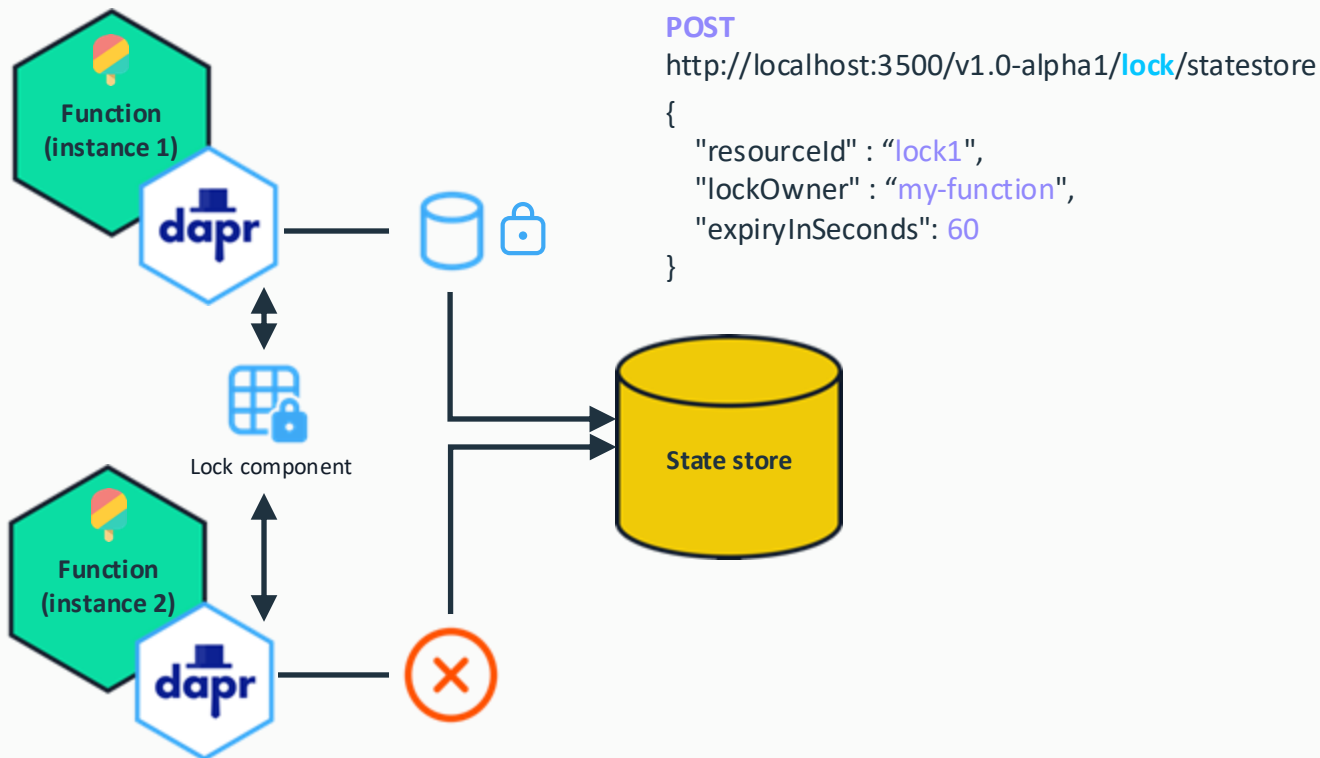
**The distributed lock API provides mutually exclusive access to resources.**

- Only a single instance of an application can hold a lock
- Locks are scoped to a Dapr app-id
- Uses a lease-based locking mechanism to prevent deadlocks

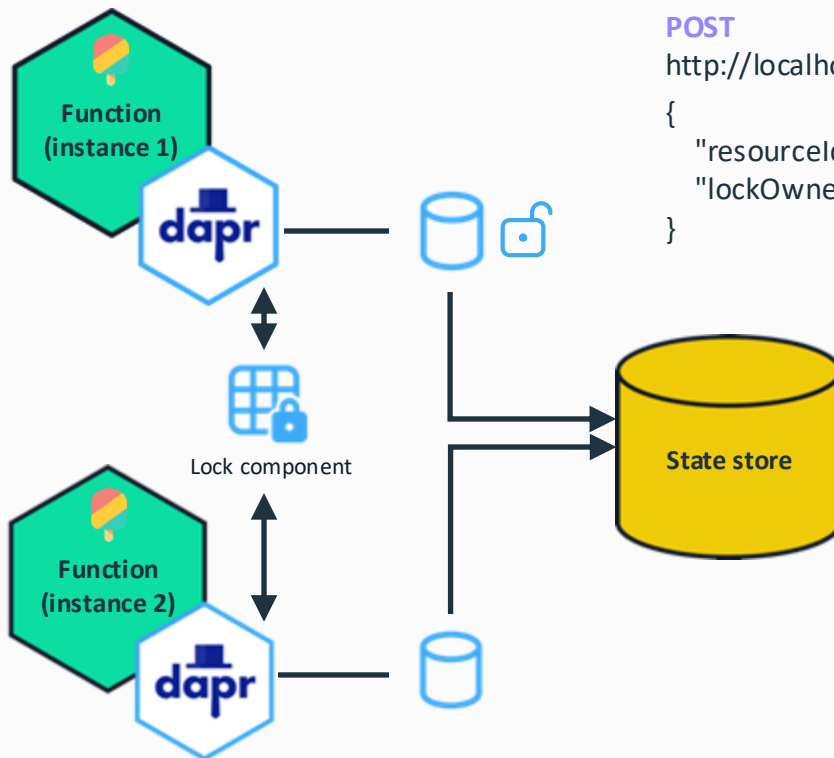
# Distributed Lock



# Distributed Lock



# Distributed Lock



POST

<http://localhost:3500/v1.0-alpha1/unlock/statestore>

```
{  
  "resourceId": "lock1",  
  "lockOwner": "my-function"  
}
```

# Service Invocation



**The service invocation API allows  
synchronous communication between  
services.**

- Service discovery via name resolution components
- Invoke HTTP and gRPC services consistently
- Configurable resiliency policies
- Built-in distributed tracing & metrics
- Access control policies & mTLS
- Chain pluggable middleware components